

---

# CS771 Major Assignment Report

---

## Group: 6 Nearest Neighbours (MLR-40)

Narayan Aniruddh Somaia (231019), Shrey Jigneshbhai Solanki (231017),  
Ayush Yadav (230272), Harshita Awasthi (230463),  
Shelly Singhal (230963), Jaini (230494)

Indian Institute of Technology Kanpur

## Part 1: Derivation of the Kernel $\tilde{K}$ for the Semi-Parametric Model

We are given the semi-parametric model

$$y = \mathbf{w}(\mathbf{z}) \cdot x + b = \mathbf{p}^\top \phi(\mathbf{z}) x + b,$$

where  $x \geq 0$  is the video length,  $\mathbf{z} = (z_1, z_2) \in \mathbb{R}^2$  contains popularity/difficulty values, and  $\mathbf{w}(\mathbf{z}) = \mathbf{p}^\top \phi(\mathbf{z})$  varies with  $\mathbf{z}$ . The feature map  $\phi$  corresponds to the polynomial kernel:

$$K(\mathbf{z}_1, \mathbf{z}_2) = (\mathbf{z}_1^\top \mathbf{z}_2 + c)^d.$$

Our goal is to express this semi-parametric model as a purely non-parametric kernel regression model compatible with `sklearn.kernel_ridge`, which does not include a bias term internally. To do this, we must construct a new feature map  $\psi$  and a corresponding kernel  $\tilde{K}$  such that

$$\tilde{\mathbf{p}}^\top \psi(x, \mathbf{z}) = \mathbf{p}^\top \phi(\mathbf{z}) x + b.$$

### Motivation for the Augmented Feature Map

Since the slope term  $\mathbf{p}^\top \phi(\mathbf{z})$  must scale with  $x$ , the feature map must contain  $x\phi(\mathbf{z})$ . Also, because kernel ridge regression has *no explicit bias*, the constant  $b$  has to be absorbed into the feature map via a constant coordinate.

This naturally leads us to define

$$\psi(x, \mathbf{z}) = \begin{bmatrix} x\phi(\mathbf{z}) \\ 1 \end{bmatrix}, \quad \tilde{\mathbf{p}} = \begin{bmatrix} \mathbf{p} \\ b \end{bmatrix}.$$

Then

$$\tilde{\mathbf{p}}^\top \psi(x, \mathbf{z}) = \mathbf{p}^\top (x\phi(\mathbf{z})) + b = \mathbf{p}^\top \phi(\mathbf{z}) \cdot x + b,$$

which matches the original model.

### Kernel Computation

For two inputs  $A = (x_1, \mathbf{z}_1)$  and  $B = (x_2, \mathbf{z}_2)$ ,

$$\tilde{K}(A, B) = \psi(x_1, \mathbf{z}_1)^\top \psi(x_2, \mathbf{z}_2) = x_1 x_2 \phi(\mathbf{z}_1)^\top \phi(\mathbf{z}_2) + 1.$$

Using the polynomial kernel identity,

$$\phi(\mathbf{z}_1)^\top \phi(\mathbf{z}_2) = (\mathbf{z}_1^\top \mathbf{z}_2 + c)^d,$$

we finally obtain:

$$\boxed{\tilde{K}((x_1, \mathbf{z}_1), (x_2, \mathbf{z}_2)) = x_1 x_2 (\mathbf{z}_1^\top \mathbf{z}_2 + c)^d + 1.}$$

This reduces the semi-parametric regression problem to a standard kernel regression problem.

## Part 2: Hyperparameter Selection for the Semi–Parametric Kernel Model

In this part, we tune the hyperparameters of the polynomial kernel used inside our semi–parametric kernel:

$$d \in \{1, 2, 3, 4, 5, 6\}, \quad c \in \{0, 0.5, 1, 1.5, 2, 5\}.$$

Predictive performance is measured using the  $R^2$  score on the test set:

$$R^2 = 1 - \frac{\sum(y_i - \hat{y}_i)^2}{\sum(y_i - \bar{y})^2}.$$

### Step 1: Hyperparameter Grid Search

For each pair  $(d, c)$  in the grid, the following steps were carried out:

1. Construct the **training Gram matrix**  $G_{\text{train}}$  using the semi-parametric kernel:

$$K((x_i, z_i), (x_j, z_j)) = x_i x_j (z_i^\top z_j + c)^d + 1.$$

2. Train a **KernelRidge** regressor with the precomputed kernel  $G_{\text{train}}$ .
3. Compute the  $R^2$  score on the test set using the corresponding **test Gram matrix**  $G_{\text{test}}$  and the `score` method of the regressor.
4. Repeat the above steps  $n_{\text{trials}}$  times to obtain stable estimates, then average the  $R^2$  score and computation times (kernel + training).

We evaluated all 36 combinations of  $(d, c)$ .  
Table 1 lists the obtained test scores:

Table 1: Test  $R^2$  scores and average computation time for different  $(d, c)$  values.

Degree $d$	Coefficient $c$	Test $R^2$ Score	Avg Time (s)
1	0	0.78984	0.55
2	0	0.57140	0.56
3	0	0.44039	0.56
4	0	0.36277	0.54
5	0	0.30730	0.55
6	0	0.26561	0.52
1	0.5	0.96969	0.54
2	0.5	0.96992	0.55
3	0.5	0.96994	0.57
4	0.5	0.96990	0.55
5	0.5	0.96979	0.55
6	0.5	0.96971	0.52
1	1	0.96969	0.55
2	1	0.96992	0.52
3	1	0.96993	0.55
4	1	0.96986	0.59
5	1	0.96970	0.60
6	1	0.96958	0.59
1	1.5	0.96969	0.61
2	1.5	0.96991	0.60
3	1.5	0.96993	0.58
4	1.5	0.96984	0.60
5	1.5	0.96969	0.60
6	1.5	0.96957	0.60
1	2	0.96969	0.60
2	2	0.96991	0.61
3	2	0.96993	0.59
4	2	0.96983	0.59
5	2	0.96968	0.58
6	2	0.96956	0.56
1	5	0.96968	0.58
2	5	0.96991	0.60
3	5	0.96993	0.59
4	5	0.96982	0.57
5	5	0.96967	0.59
6	5	0.96954	0.58

A heatmap visualization of test  $R^2$  scores is shown in Fig. 1. The numbers in parentheses indicate the *average computation time (kernel + training) in seconds* for each  $(d, c)$  pair.

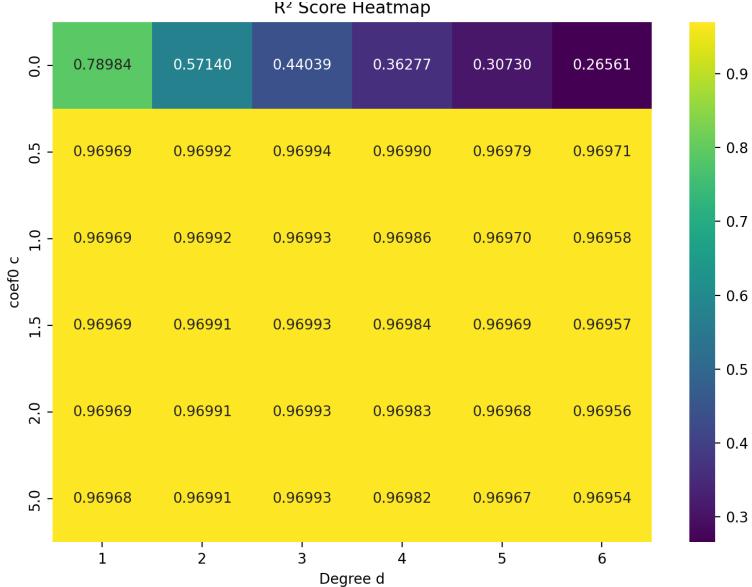


Figure 1: Heatmap of test  $R^2$  scores for various  $(d, c)$  values. Average time (s) is shown in parentheses below each  $R^2$ .

## Step 2: Model Selection

From the grid search, the highest  $R^2$  score is achieved at  $(d = 3, c = 0.5)$ . However, this score is only slightly better than the score for  $(d = 2, c = 1)$ :

$$R_{\text{gain}}^2 < 2 \times 10^{-5},$$

which is practically negligible.

At the same time, using  $d = 3$  makes the kernel more complex, which increases the computation time for both the kernel and model training. Since  $(d = 2, c = 1)$  gives almost the same accuracy but requires less computation, it is the better choice in practice.

Therefore, we select the **computationally efficient** configuration:

$$d^* = 2, \quad c^* = 1$$

with the corresponding test performance:

$$R_{\text{test}}^2 = 0.96992.$$

In short, we are prioritizing a combination that achieves nearly the best accuracy while keeping the model fast and simple.

## Discussion

- Very small  $c$  leads to underfitting (poor performance when  $c = 0$ ).
- Kernels with  $d > 3$  show no significant improvement but higher cost.
- $(d = 2, c = 1)$  achieves strong accuracy while minimizing model complexity.

Thus, this choice offers the best trade-off between predictive performance and efficiency, making it well suited for later parts of the assignment.

## Part 4: Method for Inverting the XOR Arbiter PUF

We are given a linear model

$$\mathbf{w} \in \mathbb{R}^{1089},$$

corresponding to a 2-XOR Arbiter PUF built from two independent 32-bit Arbiter PUFs. Since  $(32 + 1)^2 = 1089$ , the model can naturally be reshaped into a  $33 \times 33$  matrix

$$W = \text{reshape}(\mathbf{w}, 33, 33).$$

The goal is not to recover the original physical delays exactly (this is impossible due to scale and shift symmetries), but rather to construct *some* valid set of non-negative delays that generates a linear model identical to  $\mathbf{w}$ , as required in the assignment.

To describe our inversion method clearly, we first recall how the forward mapping from delays to the model vector is defined.

### Forward Model of a Single Arbiter PUF

A 32-bit Arbiter PUF is described by four delay parameters per stage:

$$(p_i, q_i, r_i, s_i), \quad 0 \leq i \leq 31.$$

Following the derivation from class, we define the delay-difference parameters

$$\alpha_i = \frac{p_i - q_i + r_i - s_i}{2}, \quad \beta_i = \frac{p_i - q_i - r_i + s_i}{2}.$$

These quantities determine the PUF's  $(32 + 1)$ -dimensional linear model  $u = (u_0, \dots, u_{32})^\top$  via the recurrence

$$u_0 = \alpha_0, \quad u_i = \alpha_i + \beta_{i-1} \quad (1 \leq i \leq 31), \quad u_{32} = \beta_{31}.$$

Thus a single Arbiter PUF induces a linear function

$$f(C) = \text{sign}(u^\top \Phi(C)),$$

where  $\Phi(C)$  is the standard parity-preprocessed challenge vector.

For this assignment, we focus only on the linear model  $u$  and not on  $f(C)$  itself.

### Forward Model of a 2-XOR Arbiter PUF

Let  $u, v \in \mathbb{R}^{33}$  denote the linear models of the two individual Arbiter PUFs. The XOR of their responses produces a combined model

$$w = u \otimes v,$$

where  $\otimes$  denotes the Kronecker product. Concretely, if  $u = (u_0, \dots, u_{32})^\top$  and  $v = (v_0, \dots, v_{32})^\top$ , then

$$w_{(i,j)} = u_i v_j, \quad 0 \leq i, j \leq 32.$$

Equivalently, the 1089-dimensional vector  $\mathbf{w}$  can be written as

$$\mathbf{w} = \text{vec}(uv^\top),$$

so that reshaping gives

$$W = uv^\top.$$

In practice, numerical noise and non-idealities mean that  $W$  is only approximately rank-1, but we still have

$$\text{rank}(W) \approx 1,$$

and this near rank-1 structure is what we exploit to invert the model.

### Why the First Row and First Column Encode $u$ and $v$

If  $W = uv^\top$  held exactly, then

$$W_{0,j} = u_0 v_j, \quad W_{i,0} = u_i v_0.$$

In other words, the first row is proportional to  $v^\top$  and the first column to  $u$ . Even when  $W$  is only approximately rank-1, these directions remain good summaries of the underlying models.

Motivated by this, we extract

$$r = W_{0,:}, \quad c = \frac{W_{:,0}}{r_0},$$

which (up to an overall scale) serve as surrogates for  $v^\top$  and  $u$  respectively. The normalization by  $r_0$  compensates for the scalar ambiguity

$$u \otimes v = \frac{1}{\gamma}(\gamma u) \otimes v,$$

and ensures that  $r$  and  $c$  are on a consistent scale.

We can now treat  $r$  and  $c$  as approximate Arbiter-PUF model vectors and invert them stage-wise.

### Inverting a Single Arbiter-PUF Model

Consider one of the extracted vectors, say

$$v = (v_0, \dots, v_{31}, v_{32})^\top.$$

For an ideal Arbiter PUF, the recurrence described earlier can in principle be inverted:

$$\alpha_0 = v_0, \quad \alpha_i = v_i - \beta_{i-1}, \quad \beta_{31} = v_{32}.$$

However, the pairs  $(\alpha_i, \beta_i)$  do not uniquely determine  $(p_i, q_i, r_i, s_i)$ , because many different non-negative delay assignments lead to the same  $\alpha_i$  and  $\beta_i$ . The assignment explicitly notes that uniqueness is *not* required: any consistent set of non-negative delays is acceptable.

To construct such delays efficiently, we interpret the first 32 coordinates of  $v$  as a baseline pattern of stagewise delay differences,

$$x = (v_0, \dots, v_{31}),$$

and the last coordinate  $v_{32}$  as a terminal offset that mainly influences the final stage:

$$y = (0, \dots, 0, v_{32})^\top.$$

We then form two simple candidates:

$$s = x + y, \quad d = x - y.$$

These play the role of plausible stagewise differences,

$$X_i = p_i - q_i, \quad Y_i = r_i - s_i.$$

We apply this construction independently to both  $r$  and  $c$ , giving us four candidate difference vectors in total.

### Recovering Non-Negative Delays

Each difference vector  $z$  is converted into valid non-negative delays using

$$p_i = \max(z_i, 0), \quad q_i = \max(-z_i, 0),$$

which ensures

$$z_i = p_i - q_i$$

with  $p_i, q_i \geq 0$ . Applying the same construction to another difference vector yields  $(r_i, s_i)$ .

Thus each  $z$  gives us two delay vectors, and using both  $s$  and  $d$  for each of the two model directions  $(r, c)$  produces the required 8 non-negative delay vectors.

## Summary

Our complete inversion procedure can be summarized as follows:

1. Use the Kronecker-product structure to reshape the XOR-PUF model into  $W = \text{reshape}(w, 33, 33)$ .
2. Extract model directions from the first row and normalized first column, which approximate  $v^\top$  and  $u$ .
3. For each extracted model, form two stagewise difference candidates  $s$  and  $d$ , capturing different ways the terminal offset can influence the stages.
4. Convert each difference vector into valid non-negative delays via

$$p_i = \max(z_i, 0), \quad q_i = \max(-z_i, 0).$$

## References

- Course slides and lecture notes by Purushottam Kar, CS771, IIT Kanpur.